

Clean-run observations

10 real LLM runs · 10 unique evaluated task(s) · descriptive records

Run	Model	Task result passed/evaluated	LLM requests	Tool calls /errors	Tokens input/output	Time (s)	Event log
R1	gpt-oss-120b	1/1	3	2/1	5,441/252	132.678	24 events
R2	gpt-oss-120b	1/1	2	1/0	4,216/110	86.602	15 events
R3	gpt-oss-120b	1/1	3	2/0	6,489/300	177.753	25 events
R4	gpt-oss-120b	1/1	4	3/0	8,135/283	261.179	35 events
R5	gpt-oss-120b	1/1	2	1/0	3,864/113	130.333	15 events
R6	gpt-oss-120b	1/1	3	2/0	7,652/491	196.514	25 events
R7	gpt-oss-120b	1/1	4	3/0	8,787/233	261.295	36 events
R8	gpt-oss-120b	1/1	4	3/1	9,032/582	261.624	35 events
R9	gpt-oss-120b	1/1	3	2/0	5,821/214	195.275	25 events
R10	gpt-oss-120b	1/1	3	2/0	5,840/318	195.492	25 events

Costs are per run. Repeated runs of one task are not independent task samples.

* Recording metadata absent. Times are observations, not controlled overhead measurements.