

Clean-run observations

3 real LLM runs · 1 unique evaluated task(s) · descriptive records

Run	Model	Task result passed/evaluated	LLM requests	Tool calls /errors	Tokens input/output	Time (s)	Event log
R1	gpt-oss-120b	1/1	3	2/1	5,441/196	2.015	Unrecorded*
R2	gpt-oss-120b	1/1	3	2/1	5,441/196	2.327	Unrecorded*
R3	gpt-oss-120b	1/1	3	2/1	5,441/196	2.304	24 events

Costs are per run. Repeated runs of one task are not independent task samples.

* Recording metadata absent. Times are observations, not controlled overhead measurements.