

# Clean-run observations

3 real LLM runs · 1 unique evaluated task(s) · descriptive records

| Run | Model        | Task result<br>passed/evaluated | LLM<br>requests | Tool calls<br>/errors | Tokens<br>input/output | Time<br>(s) | Event log    |
|-----|--------------|---------------------------------|-----------------|-----------------------|------------------------|-------------|--------------|
| R1  | gpt-oss-120b | 1/1                             | 3               | 2/1                   | 5,441/196              | 2.015       | No metadata* |
| R2  | gpt-oss-120b | 1/1                             | 3               | 2/1                   | 5,441/196              | 2.327       | No metadata* |
| R3  | gpt-oss-120b | 1/1                             | 3               | 2/1                   | 5,441/196              | 2.304       | 24 events    |

Costs are per run. Repeated runs of one task are not independent task samples.

\* Recording metadata absent. Times are observations, not controlled overhead measurements.